


DaQuaMRec @ RecSys 2026 Workshop

Multimodal LLMs in Action: Addressing Cold-Start Recommendation at Xiaohongshu, a Leading Chinese Social Media Platform

Guillaume Salha-Galvan

Shanghai Jiao Tong University (SJTU)
SJTU Paris Elite Institute of Technology

Minneapolis, MN · September 28, 2026



SPEAKER INTRODUCTION

Guillaume Salha-Galvan

ASSOCIATE PROFESSOR IN AI, SHANGHAI JIAO TONG UNIVERSITY

- Faculty member, **SJTU Paris Elite Institute of Technology**
- Founder and Principal Investigator, **Sino-French Data and AI Lab** (15 members)
- Ph.D. in Computer Science, **École Polytechnique**, France

ACTIVE MEMBER, ACM RECSYS COMMUNITY

- Co-author of **10 accepted RecSys papers** since 2020, across multiple tracks
- **3 papers** shortlisted for best paper awards (2020, 2021, 2024)
- Regular **Program Committee member** and workshop organizer

BRIDGING ACADEMIA AND INDUSTRY

- Nearly **10 years of industry experience** (Deezer, Kibo Ryoku)
- Research contributions powering **recommender systems used by millions of users**



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY



巴黎卓越工程师学院
SJTU Paris Elite Institute of Technology

COMPANY INTRODUCTION

Xiaohongshu



COMPANY INTRODUCTION

Xiaohongshu (RedNote) — The “Chinese Instagram”



Large-Scale CTR-Based Recommendation at Xiaohongshu

XIAOHONGSHU (REDNOTE)


300M+

monthly active users

Click-Through Rate (CTR) prediction models **rank** items for each user on the **Explore Feed**:

- **Content Feed**: personalized posts
- **Display Ads**: personalized ads

They rely on **item ID embeddings**, which capture collaborative signals but require sufficient interaction history



Large-Scale CTR-Based Recommendation at Xiaohongshu

XIAOHONGSHU (REDNOTE)

300M+

monthly active users

Click-Through Rate (CTR) prediction models **rank** items for each user on the **Explore Feed**:

- **Content Feed**: personalized posts
- **Display Ads**: personalized ads

They rely on **item ID embeddings**, which capture collaborative signals but require sufficient interaction history

The item cold-start problem

New posts are uploaded with little or no interaction history, leaving their ID embeddings poorly estimated

- **New items must be served immediately**
Waiting for interaction history would degrade user experience
- **Volume is massive**
New items arrive at massive scale, amplifying the issue
- **Production compatibility is critical**
Cold-start solutions should reuse mature production CTR architectures already deployed at scale

Cold-Start CTR Prediction

CTR MODEL f_θ

Predicts the probability that user u clicks on item i , given:

- A user embedding $e_u \in \mathbb{R}^d$
- An **item ID embedding** $e_i \in \mathbb{R}^d$
- Contextual features x_{ui}

CTR predictions are then used to rank candidate items

$$\hat{y}_{ui} = f_\theta(\mathbf{e}_u, \mathbf{e}_i, \mathbf{x}_{ui}) \in [0, 1]$$

Cold-Start CTR Prediction

CTR MODEL f_θ

Predicts the probability that user u clicks on item i , given:

- A user embedding $e_u \in \mathbb{R}^d$
- An **item ID embedding** $e_i \in \mathbb{R}^d$
- Contextual features x_{ui}

CTR predictions are then used to rank candidate items

$$\hat{y}_{ui} = f_\theta(\mathbf{e}_u, \mathbf{e}_i, \mathbf{x}_{ui}) \in [0, 1]$$

COLD-START CTR

New items/posts **lack sufficient interaction history**:

- Poorly estimated ID embeddings
- Degraded CTR estimation and ranking quality
- But items are associated with **multimodal content (images, text...)**

Goal: use content to learn a **proxy embedding** $p_i \in \mathbb{R}^d$

$$\hat{y}_{ui} = f_\theta(\mathbf{e}_u, \mathbf{p}_i, \mathbf{x}_{ui}) \in [0, 1]$$

PROBLEM FORMULATION

Cold-Start CTR Prediction

CTR MODEL f_θ

Predicts the probability that user u clicks on item i , given:

- A user embedding $e_u \in \mathbb{R}^d$
- An **item ID embedding** $e_i \in \mathbb{R}^d$
- Contextual features x_{ui}

CTR predictions are then used to rank candidate items

$$\hat{y}_{ui} = f_\theta(\mathbf{e}_u, \mathbf{e}_i, \mathbf{x}_{ui}) \in [0, 1]$$

COLD-START CTR

New items/posts **lack sufficient interaction history**:

- Poorly estimated ID embeddings
- Degraded CTR estimation and ranking quality
- But items are associated with **multimodal content** (images, text...)

Goal: use content to learn a **proxy embedding** $p_i \in \mathbb{R}^d$

$$\hat{y}_{ui} = f_\theta(\mathbf{e}_u, \mathbf{p}_i, \mathbf{x}_{ui}) \in [0, 1]$$

THE FUNDAMENTAL QUESTION

Can we align multimodal representations with ID embeddings while maximally reusing the existing CTR ranker?


Cold-Start CTR Prediction

$$\hat{y}_{ui} = f_{\theta}(\mathbf{e}_u, \mathbf{e}_i, \mathbf{x}_{ui}) \in [0, 1]$$

Goal: use content to learn a **proxy embedding** $p_i \in \mathbb{R}^d$

$$\hat{y}_{ui} = f_{\theta}(\mathbf{e}_u, \mathbf{p}_i, \mathbf{x}_{ui}) \in \mathbb{R}$$


Can we align multimodal representations with ID embeddings while maximally retaining information?



MOTIVATION

Why Is Aligning Content with Collaborative ID Embeddings Hard at Xiaohongshu?

MovieLens-1M ID Embeddings (SASRec)



Direct MLP Content-to-ID Mapping


Examples: CB2CF, CLCRec, GoRec, GAR

- Works when ID space is **content-clustered**
Ex: academic benchmarks
- Struggles with **irregular ID spaces**


MOTIVATION

Why Is Aligning Content with Collaborative ID Embeddings Hard at Xiaohongshu?

MovieLens-1M ID Embeddings (SASRec)



Production ID Embeddings from Xiaohongshu



Direct MLP Content-to-ID Mapping

Examples: CB2CF, CLCRec, GoRec, GAR

- Works when ID space is **content-clustered**
Ex: *academic benchmarks*
- Struggles with **irregular ID spaces**

Irregular

Non-Content-Clustered

MOTIVATION

Why Is Aligning Content with Collaborative ID Embeddings Hard at Xiaohongshu?


Reco-Aware Multimodal Embeddings (QARM)



MovieLens-1M ID Embeddings (SASRec)



Production ID Embeddings from Xiaohongshu



Co-Occurrence Alignment

Examples: QARM, MOON, NoteLLM2

- Pull co-occurring items close in content space
- Often yields **clustered semantic geometry**
- Useful; **not a drop-in proxy** for ID embeddings

Direct MLP Content-to-ID Mapping

Examples: CB2CF, CLCRec, GoRec, GAR

- Works when ID space is **content-clustered**
Ex: *academic benchmarks*
- Struggles with **irregular ID spaces**

Irregular

Non-Content-Clustered

MOTIVATION

Why Is Aligning Content with Collaborative ID Embeddings Hard at Xiaohongshu?


Reco-Aware Multimodal Embeddings (QARM)



MovieLens-1M ID Embeddings (SASRec)



Production ID Embeddings from Xiaohongshu



Co-Occurrence Alignment

Examples: QARM, MOON, NoteLLM2

- Pull co-occurring items close in content space
- Often yields **clustered semantic geometry**
- Useful; **not a drop-in proxy** for ID embeddings

Direct MLP Content-to-ID Mapping

Examples: CB2CF, CLCRec, GoRec, GAR

- Works when ID space is **content-clustered**
Ex: *academic benchmarks*
- Struggles with **irregular ID spaces**

Irregular

Non-Content-Clustered

Common limitation: neither is fully aligned with the **production CTR objective** in general

IDProxy

**A Two-Stage, Coarse-to-Fine, End-to-End Framework
based on Multimodal LLMs**

IDProxy: A Two-Stage, Coarse-to-Fine, End-to-End Framework

1

COARSE PROXY GENERATION

Stage 1 · Offline Contrastive Alignment

- **Multimodal LLM (MLLM)** encodes each item's **images + text**
- Learn **proxy** p_i^{coarse} in the ID embedding space
- **Contrastive loss** against ℓ_2 -normalized ID embeddings
- Filter low-interactions items to get reliable training targets

2

CTR-AWARE FINE PROXY ALIGNMENT

Stage 2 · End-to-End Training with the CTR Ranker

- Multi-layer **MLLM hidden states** → richer semantics
- Lightweight adaptor + residual gating
- **Refined proxy** p_i^{fine} enters feature-cross + target attention
- MLLM frozen; adaptor + ranker **trained with CTR loss**

Reuse the ranker's architecture and structural priors, with no separate pipeline or hand-crafted alignment

IDProxy: A Two-Stage, Coarse-to-Fine, End-to-End Framework

1

COARSE PROXY GENERATION

Stage 1 · Offline Contrastive Alignment

- **Multimodal LLM (MLLM)** encodes each item's **images + text**
- Learn **proxy** p_i^{coarse} in the ID embedding space
- **Contrastive loss** against ℓ_2 -normalized ID embeddings
- Filter low-interactions items to get reliable training targets

2

CTR-AWARE FINE PROXY ALIGNMENT

Stage 2 · End-to-End Training with the CTR Ranker

- Multi-layer **MLLM hidden states** → richer semantics
- Lightweight adaptor + residual gating
- **Refined proxy** p_i^{fine} enters feature-cross + target attention
- MLLM frozen; adaptor + ranker **trained with CTR loss**

Reuse the ranker's architecture and structural priors, with no separate pipeline or hand-crafted alignment

Stage 1 · MLLM-Based Coarse Proxy Generation



Stage 1 · MLLM-Based Coarse Proxy Generation

Multimodal Content



Note title: 2026 Beijing Chinese
Program Language
Note content: [...]

↓ for “warm” items, prompt template with [EMB] token

Multimodal LLM (InternVL, offline)
Final Hidden States H_i

↓ attention pool $g(\cdot) \rightarrow$ MLP $\varphi \rightarrow$ ℓ_2 norm

$\tilde{h}_i \rightarrow$ Coarse Proxy $\mathbf{p}_i^{\text{coarse}} = \tilde{\mathbf{h}}_i \in \mathbb{R}^d$

Stage 1 · MLLM-Based Coarse Proxy Generation

Multimodal Content



Note title: 2026 Beijing Chinese
Program Language
Note content: [...]

↓ for “warm” items, prompt template with [EMB] token

Multimodal LLM (InternVL, offline)
Final Hidden States H_i

↓ attention pool $g(\cdot) \rightarrow$ MLP $\varphi \rightarrow$ ℓ_2 norm

$\tilde{h}_i \rightarrow$ Coarse Proxy $\mathbf{p}_i^{\text{coarse}} = \tilde{\mathbf{h}}_i \in \mathbb{R}^d$

PROXY ALIGNMENT LOSS

$$\mathcal{L}_{\text{PAL}} = -\frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \log \frac{\exp(\tilde{\mathbf{h}}_i^\top \mathbf{e}_i / \tau_c)}{\sum_{j \in \mathcal{B}} \exp(\tilde{\mathbf{h}}_i^\top \mathbf{e}_j / \tau_c)}$$

INTUITION

Contrastive objective pulls each MLLM-derived embedding toward its own **ID embedding** and pushes it away from other items in the batch

Result: an MLLM coarsely aligned with the production ID space

Stage 1 · MLLM-Based Coarse Proxy Generation

Multimodal Content



Note title: 2026 Beijing Chinese
Program Language
Note content: [...]

↓ for “warm” items, prompt template with [EMB] token

Multimodal LLM (InternVL, offline)
Final Hidden States H_i

↓ attention pool $g(\cdot) \rightarrow$ MLP $\varphi \rightarrow$ ℓ_2 norm

$\tilde{h}_i \rightarrow$ Coarse Proxy $\mathbf{p}_i^{\text{coarse}} = \tilde{\mathbf{h}}_i \in \mathbb{R}^d$

PROXY ALIGNMENT LOSS

$$\mathcal{L}_{\text{PAL}} = -\frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \log \frac{\exp(\tilde{\mathbf{h}}_i^\top \mathbf{e}_i / \tau_c)}{\sum_{j \in \mathcal{B}} \exp(\tilde{\mathbf{h}}_i^\top \mathbf{e}_j / \tau_c)}$$

INTUITION

Contrastive objective pulls each MLLM-derived embedding toward its own **ID embedding** and pushes it away from other items in the batch


Result: an MLLM coarsely aligned with the production ID space

But coarse alignment alone cannot exploit the CTR ranker's end-to-end structure. Enter Stage 2.

Stage 2 · CTR-Aware Fine Proxy Alignment

HIERARCHICAL REPRESENTATIONS

Extract **3 subgroups** from **MLLM hidden layer representations**:




Learn **hierarchical semantic representations** using Stage 1's attention pooling g :

$$\mathbf{z}_i^{(l)} = g(\mathbf{H}_i^{(l)}) \in \mathbb{R}^D, \text{ for } l \in \{l_{n1}, l_{n2}, l_{n3}\}.$$

Stage 2 · CTR-Aware Fine Proxy Alignment

HIERARCHICAL REPRESENTATIONS

Extract **3 subgroups** from **MLLM hidden layer representations**:



Learn **hierarchical semantic representations** using Stage 1's attention pooling g :

$$\mathbf{z}_i^{(l)} = g(\mathbf{H}_i^{(l)}) \in \mathbb{R}^D, \text{ for } l \in \{l_{n1}, l_{n2}, l_{n3}\}.$$


MOTIVATION

Research shows that **intermediate MLLM representations** can capture richer, more transferable semantics (*see paper*)

Stage 2 · CTR-Aware Fine Proxy Alignment

HIERARCHICAL REPRESENTATIONS

Extract **3 subgroups** from **MLLM hidden layer representations**:



Learn **hierarchical semantic representations** using Stage 1's attention pooling g :

$$\mathbf{z}_i^{(l)} = g(\mathbf{H}_i^{(l)}) \in \mathbb{R}^D, \text{ for } l \in \{l_{n1}, l_{n2}, l_{n3}\}.$$

LIGHTWEIGHT PROXY REFINEMENT

Concatenation + **lightweight MLP** model with fewer parameters than the MLLM:

$$\mathbf{p}_i^{\text{fine}} = \tilde{\phi}(\text{Concat}(\mathbf{z}_i^{(l_{n1})}, \mathbf{z}_i^{(l_{n2})}, \mathbf{z}_i^{(l_{n3})})) \in \mathbb{R}^d$$

In practice: add **residual gating** modules to limit overlap between *coarse* and *fine*


MOTIVATION

Research shows that **intermediate MLLM representations** can capture richer, more transferable semantics (*see paper*)

Stage 2 · CTR-Aware Fine Proxy Alignment

HIERARCHICAL REPRESENTATIONS

Extract **3 subgroups** from **MLLM hidden layer representations**:



Learn **hierarchical semantic representations** using Stage 1's attention pooling g :

$$\mathbf{z}_i^{(l)} = g(\mathbf{H}_i^{(l)}) \in \mathbb{R}^D, \text{ for } l \in \{l_{n1}, l_{n2}, l_{n3}\}.$$

LIGHTWEIGHT PROXY REFINEMENT

Concatenation + **lightweight MLP** model with fewer parameters than the MLLM:

$$\mathbf{p}_i^{\text{fine}} = \tilde{\phi}(\text{Concat}(\mathbf{z}_i^{(l_{n1})}, \mathbf{z}_i^{(l_{n2})}, \mathbf{z}_i^{(l_{n3})})) \in \mathbb{R}^d$$

In practice: add **residual gating** modules to limit overlap between *coarse* and *fine*

INTEGRATION INTO CTR RANKER f_θ

Proxies as f_θ inputs, used for **feature crossing** and **target attention**:

$$\hat{y}_{ui} = f_\theta(\mathbf{e}_u, \mathbf{e}_i, \mathbf{p}_i^{\text{coarse}}, \mathbf{p}_i^{\text{fine}}, \mathbf{x}_{ui})$$

Weights are **optimized end-to-end** under **CTR objective**:

$$\mathcal{L}_{\text{CTR}} = -\frac{1}{|\mathcal{D}|} \sum_{(u, i, \mathbf{x}_{ui}, y_{ui}) \in \mathcal{D}} [y_{ui} \log \hat{y}_{ui} + (1 - y_{ui}) \log (1 - \hat{y}_{ui})]$$

Click/no-click classification, MLLM frozen

MOTIVATION

Research shows that **intermediate MLLM representations** can capture richer, more transferable semantics (*see paper*)

Designed for **large-scale deployment** on Xiaohongshu

Four Key Questions

Q1 · MLLM-BASED COARSE STAGE

How effective is the **MLLM-based coarse stage** compared with existing **alignment methods**?

Four Key Questions

Q1 · MLLM-BASED COARSE STAGE

How effective is the **MLLM-based coarse stage** compared with existing **alignment methods**?

Q2 · CTR-AWARE REFINEMENT

What is the contribution of the **fine stage**, i.e., what is the impact of **end-to-end modeling** and **structural reuse**?

Four Key Questions

Q1 · MLLM-BASED COARSE STAGE

How effective is the **MLLM-based coarse stage** compared with existing **alignment methods**?

Q2 · CTR-AWARE REFINEMENT

What is the contribution of the **fine stage**, i.e., what is the impact of **end-to-end modeling** and **structural reuse**?

Q3 · COLD-START EFFECTIVENESS

Does IDProxy effectively alleviate the **cold-start problem** for newly published content on Xiaohongshu?

Four Key Questions

Q1 · MLLM-BASED COARSE STAGE

How effective is the **MLLM-based coarse stage** compared with existing **alignment methods**?

Q2 · CTR-AWARE REFINEMENT

What is the contribution of the **fine stage**, i.e., what is the impact of **end-to-end modeling** and **structural reuse**?

Q3 · COLD-START EFFECTIVENESS

Does IDProxy effectively alleviate the **cold-start problem** for newly published content on Xiaohongshu?

Q4 · ONLINE BUSINESS IMPACT

What is the impact of IDProxy on **real-world business metrics** at Xiaohongshu?

Offline CTR Prediction (Historical Click Logs from 2025)

BASE: XIAOHONGSHU'S PRODUCTION BASELINE

A highly optimized **industrial CTR system** with ID-based **feature interactions** and **sequential user modeling**

EXPERIMENTAL RESULTS

MLLM-based ID alignment (v3) beats:

- Xiaohongshu’s baseline
- Recommendation-aware multimodal embeddings (v1)
- Simple MLP content → ID mapping (v2)

MLLM-based alignment with **CTR-aware refinement (v4, v5)** delivers the largest gains: up to **+0.14% ΔAUC**

Table 1: Offline CTR prediction gains relative to the baseline.

Model Variant	Model ID	Δ AUC
Base (Production Baseline at Xiaohongshu)	-	0
Base + Notellm2-Like Embed	v1	+0.015%
Base + Static Vector (MLP Mapping)	v2	+0.02%
Base + IDProxy (Stage 1)	v3	+0.05%
Base + IDProxy (Stage 1 + 2, w/o Structure Reuse)	v4	+0.08%
Base + IDProxy (Stage 1 + 2)	v5	+0.14%

*Practical significance at Xiaohongshu: **+0.001 absolute AUC***

Results are statistically significant at the 1% level

Online A/B Tests on Xiaohongshu's Explore Feed (2025)

TWO A/B TESTS ON MILLIONS OF USERS

IDProxy deployed on treatment groups in two settings:

- **Content Feed:** personalized posts (August 2025)
- **Display Ads:** personalized ads (March 2025)

EXPERIMENTAL RESULTS

5-day evaluation of **new notes published within 24h:**

- Roughly **2×** larger gains than **global traffic**
- Effectively improving **cold-start ranking** with limited history

Consistent **business gains** across all measured online metrics, spanning both **Content Feed** and **Display Ads**

Table 2: Online Δ AUC gains for global notes and new notes.

	Day 1	Day 2	Day 3	Day 4	Day 5
Global Notes	+0.13%	+0.15%	+0.14%	+0.12%	+0.15%
New Notes	+0.24%	+0.32%	+0.23%	+0.27%	+0.31%

Table 3: Online A/B tests on Xiaohongshu’s Explore Feed.

Scenario	Online A/B Test Metric			
Using IDProxy in Content Feed	Time Spent	Reads	Engagements	-
	+0.22%	+0.39%	+0.5%	-
Using IDProxy in Display Ads	Impression	ADV	COST	CTR
	+1.28%	+1.93%	+1.73%	+0.23%

*Practical significance at Xiaohongshu: **+0.001 absolute AUC***

Results are statistically significant at the 1% level

Conclusion and Future Work


CONCLUSION

IDProxy was **fully deployed** in late 2025:

- The system now powers both Content Feed and Display Ads
- It serves **hundreds of millions of users daily**

Beyond its industrial impact, our study demonstrates:

- Benefits of **MLLM-based** multimodal-to-ID **alignment**
- Importance of **end-to-end refinement** with core **CTR rankers**
- Potential of **MLLM hidden states** for content representations



Conclusion and Future Work

LIMITATIONS AND FUTURE WORK

Alignment relies on contrastive learning with in-batch negatives

- Explore richer distributional **alignment objectives**?

The MLLM remains **frozen** during CTR training:

- Allow **CTR gradients** to directly **refine MLLM** representations?

Our results also motivate future research on **MLLM hidden representations** and their integration into recommender systems

IDProxy: CTR Prediction with Multimodal LLMs for Cold-Start Recommendation at Xiaohongshu

Yubin Zhang
Xiaohongshu Inc.
Shanghai, China
zhangyubin@xiaohongshu.com

Haiming Xu
Xiaohongshu Inc.
Shanghai, China
xuhaiming@xiaohongshu.com

Guillaume Salha-Galvan
Shanghai Jiao Tong University
Shanghai, China
gsalhagalvan@sjtu.edu.cn

Ruiyan Han
Fudan University
Shanghai, China
ryhan25@m.fudan.edu.cn

Feiyang Xiao
Fudan University
Shanghai, China
fyxiao24@m.fudan.edu.cn

Yanhua Huang
Xiaohongshu Inc.
Shanghai, China
yanhuahuang@xiaohongshu.com

Li Lin
Xiaohongshu Inc.
Shanghai, China
lufei2@xiaohongshu.com

Yang Luo
Xiaohongshu Inc.
Beijing, China
yangluo1@xiaohongshu.com

Yao Hu
Xiaohongshu Inc.
Beijing, China
xiahou@xiaohongshu.com

LEAD AUTHOR

Yubin ZHANG

zhangyubin@xiaohongshu.com



Scan for paper

LOOKING AHEAD

See You on Wednesday!



**Industry Track
Paper (Oral)**

**Session 5A: Industrial Ranking, CTR &
Scalable Models**

**Wednesday, September 30
2:00 – 3:30PM**

IDProxy: CTR Prediction with Multimodal LLMs for Cold-Start Recommendation at Xiaohongshu

Yubin Zhang
Xiaohongshu Inc.
Shanghai, China
zhangyubin@xiaohongshu.com

Haiming Xu
Xiaohongshu Inc.
Shanghai, China
xuhaiming@xiaohongshu.com

Guillaume Salha-Galvan
Shanghai Jiao Tong University
Shanghai, China
gsalhagalvan@sjtu.edu.cn

Ruiyan Han
Fudan University
Shanghai, China
ryhan25@m.fudan.edu.cn

Feiyang Xiao
Fudan University
Shanghai, China
fyxiao24@m.fudan.edu.cn

Yanhua Huang
Xiaohongshu Inc.
Shanghai, China
yanhuahuang@xiaohongshu.com

Li Lin
Xiaohongshu Inc.
Shanghai, China
lufei2@xiaohongshu.com

Yang Luo
Xiaohongshu Inc.
Beijing, China
yangluo1@xiaohongshu.com

Yao Hu
Xiaohongshu Inc.
Beijing, China
xiahou@xiaohongshu.com

LEAD AUTHOR

Yubin ZHANG

zhangyubin@xiaohongshu.com



Scan for paper

DaQuaMRec @ RecSys 2026 Workshop

Multimodal LLMs in Action: Addressing Cold-Start Recommendation at Xiaohongshu, a Leading Chinese Social Media Platform

Guillaume Salha-Galvan

Shanghai Jiao Tong University (SJTU)
SJTU Paris Elite Institute of Technology

Minneapolis, MN · September 28, 2026

